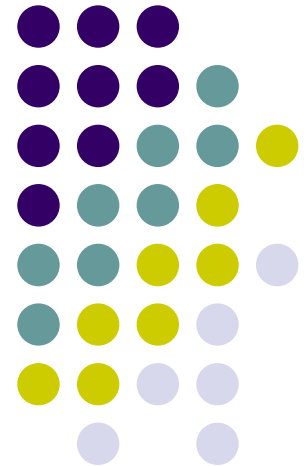


Practical Parallel Computing (実践的並列コンピューティング)

Part 1: OpenMP
No 2: Diffusion Sample
Apr 18, 2024

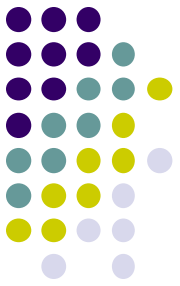
Toshio Endo
School of Computing & GSIC
endo@is.titech.ac.jp





Overview of This Course

- Part 0: Introduction
 - 2 classes
- Part 1: **OpenMP** for shared memory programming
 - 4 classes **← We are here (2/4)**
- Part 2: **GPU** programming
 - OpenACC and CUDA
 - 4 classes
- Part 3: **MPI** for distributed memory programming
 - 3~4 classes

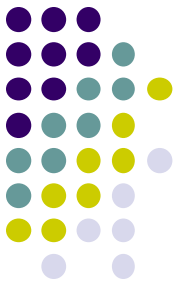


Summary of Previous Class

OpenMP is for shared-memory parallel programming

- `#pragma omp parallel` defines a parallel region, where multiple threads work simultaneously
- With `#pragma omp for`, loop-based programs can be parallelized easily
- Shared variables and private variables
- We have reviewed OpenMP version of `mm` sample


“diffusion” Sample Program



An example of diffusion phenomena:

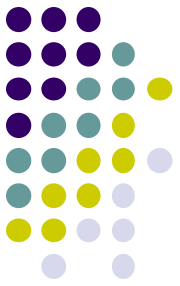
- Pour a drop of ink into a water glass

$$\frac{\partial \phi}{\partial t} = D \nabla^2 \phi(\vec{r}, t)$$



The ink spreads gradually, and finally the density becomes uniform (Figure by Prof. T. Aoki, GSIC)

- Density of ink in each point vary according to time → Simulated by computers
 - cf) Weather forecast compute wind speed, temperature, air pressure...



“diffusion” Sample on TSUBAME

Available at [/gs/bs/tga-ppcomp/24/diffusion/](https://gs/bs/tga-ppcomp/24/diffusion/)

- Execution : `./diffusion [nt]`
- nt: Number of time steps
- nx, ny: Space grid size
 - nx=20000, ny=20000 (Fixed. See the code)
 - How can we make them variables? (mm sample will be useful as a reference)
- Compute Complexity : $O(nx \times ny \times nt)$


There is also [/gs/bs/tga-ppcomp/24/diffusion-omp/](https://gs/bs/tga-ppcomp/24/diffusion-omp/)

- [NOTE] diffusion.c is not parallel
- You can use it as a start point


Expression of Space to be Simulated



- Space to be simulated are divided into grids, and expressed by arrays (2D in this sample)




- Array elements are computed via timestep, by using “previous” data





Stencil Computations

- A data point (x,y) at time $t+1$ is computed using following data
 - point (x,y) at time t
 - “Neighbor” points of (x,y) at time t



- In diffusion sample, the computation is simply “average of 5 points”
- Computations of similar type are called “**stencil computations**”
 - Frequently used in fluid simulations



Original meanings of “stencil”

Initial Conditions & Boundary Conditions




In stencil computations, following data points cannot be computed

Instead, we have to give them (for example, as input data)

- All points at $t=0$ (Initial conditions)
 - In diffusion sample, given in `init()`
- “Boundary” points for all t (Boundary conditions)
 - In diffusion sample, they are constant during simulation
→ See ranges of for-loops in `calc()`; boundaries are skipped
 - This is not good for simulation of a water glass ☹, but it’s simple...

Initial Conditions

Boundary Conditions






A Single Array Does not Work

Let us compute $t \rightarrow t+1$

- With a single 2D array (Bug! ☹️)




We need neighbor points at time t , but some have been already updated to $t+1$ ☹️


Bad new data

- With separate 2D arrays (Good 😊)

An array for t



An array for $t+1$




We can access “old” neighbor points correctly 😊



Multiple Arrays?

- We repeat update of the array for NT times



A simple way is to make arrays for all time steps

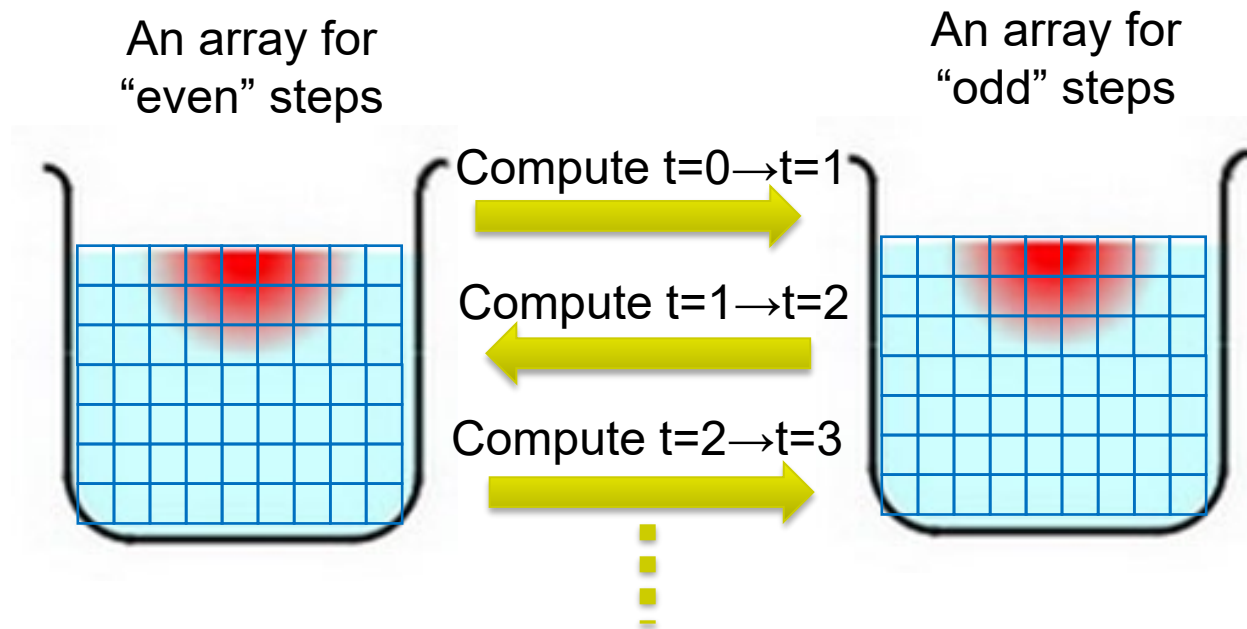
```
float data[NT+1][NY][NX]
```

- This uses too much memory
- Do we need all of $(NT+1)$ arrays?

Double Buffering Technique

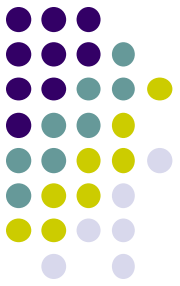


- It is sufficient to have “current” array and “next” array.
- It is better to use only “**Double buffers**”



The diffusion sample program uses
`float data[2][NY][NX];`

How We Parallelize “diffusion” (Related to Assignment [O1])



calc() takes long time, complexity is $O(n_x n_y n_t)$

It mainly uses “for” loops

→ **#pragma omp parallel for** is useful! But...

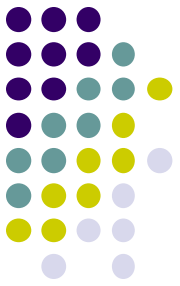
There are 3 (t, x, y) loops. Which should be parallelized?

[Hint1] Parallelizing either of spatial loop (x, y) would be good. Then spaces are divided into multiple threads

→ **[Q]** Parallelizing t loop is a not good idea. Why?

[Hint2] Take care of “pitfall in nested loops” (see slides in previous class)

Towards “Correct” Parallel Programming



There are several types of **bugs** in parallel programming

- Bugs in compile time
- Bugs in run time
 - Bugs that abort execution (cf. segmentation fault)
 - **Silent bugs → Hardest to find!**

All bugs should be avoided!



When Can We Use “omp for”?

- Loops with some (complex) forms cannot be supported, unfortunately ☹️
- The target loop must be in the following form

```
#pragma omp for
  for (i = value; i op value; incr-part)
    body
```

“*op*” : <, >, <=, >=, etc.

“*incr-part*” : i++, i--, i+=c, i-=c, etc.

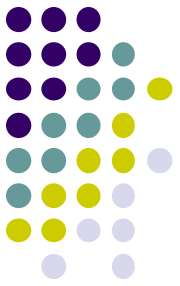
OK 😊: for (x = n; x >= 0; x-=4) ...

ERROR ☹️: for (i = 0; test(i); i++) ...

ERROR ☹️: for (p = head; p != NULL; p = p->next)

} Errors in
compile time

What are Differences between These Codes?



```
double D[100];  
:
```

Code A

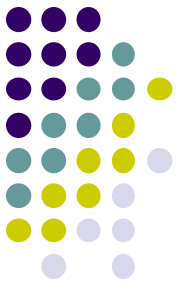
```
#pragma omp parallel for  
for (i = 0; i < 100; i++) {  
    D[i] = D[i]+1.0;  
}
```

Code B

```
#pragma omp parallel for  
for (i = 0; i < 99; i++) {  
    D[i+1] = D[i]+1.0;  
}
```


- Both codes can be compiled and executed...
- But **only code A is correct** 😊 , **code B has a bug** ☹️
 - Code B's results may be wrong

Sequential Execution and Parallel Execution of Loop



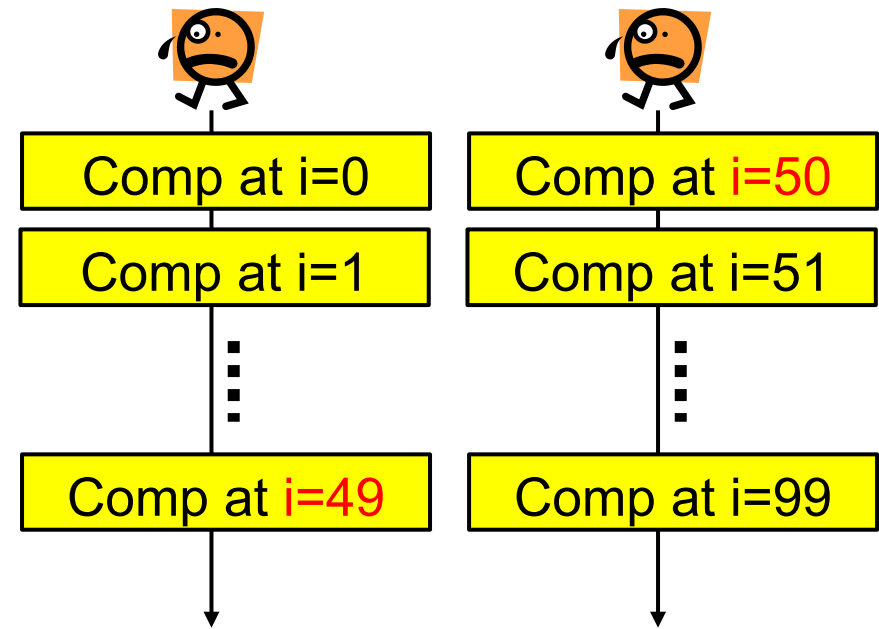
[Sequential]

for (i = 0; i < 100; i++) ...



[Parallel]

#pragma omp parallel for
for (i = 0; i < 100; i++) ...



in case of 2 threads,
i=50 is computed before i=49



Difference between Two Codes

Code A

```
#pragma omp parallel for
  for (i = 0; i < 100; i++) {
    D[i] = D[i]+1.0;
  }
```

OK

It is **ok to reorder** 100 computations

Code B

```
#pragma omp parallel for
  for (i = 0; i < 99; i++) {
    D[i+1] = D[i]+1.0;
  }
```

NG

Computations **must be done in an order** (i=0,1,2...)

➔ Parallelization breaks the order

Dependency between Computations



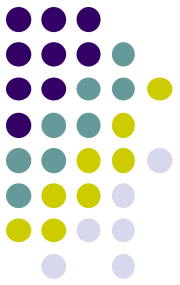
We define following sets for computation C

- Read set $R(C)$: the set of variables **read** by C
- Write set $W(C)$: the set of variables **written** by C
 - Ex) C: $x = y + z \rightarrow R(C) = \{y, z\}, W(C) = \{x\}$

We define **dependency** between C1 and C2

- If $(W(C1) \cap R(C2) \neq \emptyset)$, C1 and C2 are **dependent** (**write** vs **read**)
- If $(R(C1) \cap W(C2) \neq \emptyset)$, C1 and C2 are **dependent** (**read** vs **write**)
- If $(W(C1) \cap W(C2) \neq \emptyset)$, C1 and C2 are **dependent** (**write** vs **write**)
- Otherwise, C1 and C2 are **independent**
 - ✖ **read vs read** cases are independent

If C1 and C2 are **independent**, parallelization of C1 and C2 is safe ☺



Example of Dependency

Code A

```
#pragma omp parallel for
  for (i = 0; i < 100; i++) {
    D[i] = D[i]+1.0;    ← Ai
  }
```

$R(A_i) = \{D[i]\}$, $W(A_i) = \{D[i]\}$

All 100 computations are independent

Code B

```
#pragma omp parallel for
  for (i = 0; i < 99; i++) {
    D[i+1] = D[i]+1.0; ← Bi
  }
```

$R(B_i) = \{D[i]\}$, $W(B_i) = \{D[i+1]\}$

$R(B_{i+1}) \cap W(B_i) = \{D[i+1]\} \neq \emptyset \rightarrow \text{Dependent!}$


Dependency and Parallelism in Stencil Computations



Consider 1D stencil computation:

```
for (t = 0; t < NT; t++)
  for (x = 1; x < NX-1; x++)
     $f_{t+1,x} = (f_{t,x-1} + f_{t,x} + f_{t,x+1}) / 3.0$  /*  $c_{t,x}$  */
```


✂ This is simpler than “diffusion” (2D) sample



We let $c_{t,x}$ be computation of a single point $f_{t+1,x}$

$R(c_{t,x}) = \{f_{t,x-1}, f_{t,x}, f_{t,x+1}\}$, $W(c_{t,x}) = \{f_{t+1,x}\}$

..... $x=$ 19 20 21



✂ This figure omits double buffering technique

Discussion on Stencil: Case of Spatial Loop




```
for (t = 0; t < NT; t++)
```

```
  for (x = 1; x < NX-1; x++)    ← Is this loop parallelizable?
```

```
    ft+1,x = (ft,x-1 + ft,x + ft,x+1) / 3.0 /* ct,x */
```

- Can we compute $c_{5,20}$ and $c_{5,21}$ in parallel? (*t is same, x is different*)
 - $R(C_{5,20}) = \{f_{5,19}, f_{5,20}, f_{5,21}\}$, $W(C_{5,20}) = \{f_{6,20}\}$
 - $R(C_{5,21}) = \{f_{5,20}, f_{5,21}, f_{5,22}\}$, $W(C_{5,21}) = \{f_{6,21}\}$
 - They are **independent** 😊, for all pairs of x
 - x loop can be parallelized

Read vs. Read is Ok



Discussion on Stencil: Case of Temporal Loop



```
for (t = 0; t < NT; t++)
```

← Is this loop parallelizable?

```
  for (x = 1; x < NX-1; x++)
```

```
     $f_{t+1,x} = (f_{t,x-1} + f_{t,x} + f_{t,x+1}) / 3.0$ 
```

C_t


- Can we compute C_5 and C_6 in parallel? (t is different)

- $R(C_5) = \{f_{5,0}, \dots, f_{5,NX-1}\}$, $W(C_5) = \{f_{6,1}, \dots, f_{6,NX-2}\}$

- $R(C_6) = \{f_{6,0}, \dots, f_{6,NX-1}\}$, $W(C_6) = \{f_{7,1}, \dots, f_{7,NX-2}\}$

→ $R(C_6) \cap W(C_5) = \{f_{6,1}, \dots, f_{6,NX-2}\} \neq \emptyset$

→ They are **dependent** ☹️



dependent!!



In Assignment [O1]

- it is **OK** to parallelize x-loop or y-loop
- it is **NG** to parallelize t-loop

Advanced Topic: More Speed in Stencil

We see dependency more in detail:

- $c_{6,20}$ depends on $c_{5,19}$, $c_{5,20}$, $c_{5,21}$
 - The same point or its direct neighbor
- But not on $c_{5,22}$

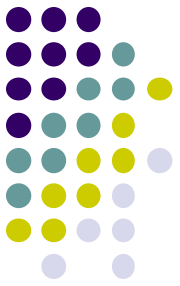


Temporal blocking technique:

After computations in (1) finish,
we can start (2)

“Trapezoids” in the same stage
can be parallelized

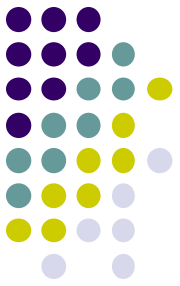
→ Speed is improved for better
access locality



Using Larger TSUBAME Resources

➔ See ppcomp-sup slides

Assignments in OpenMP Part (Abstract)



Choose one of [O1]—[O3], and submit a report
Due date: May 9 (Thu)

[O1] Parallelize “diffusion” sample program by OpenMP.

(</gs/bs/tga-ppcomp/24/diffusion/> on TSUBAME)

[O2] Parallelize “sort” sample program by OpenMP.

(</gs/bs/tga-ppcomp/24/sort/> on TSUBAME)

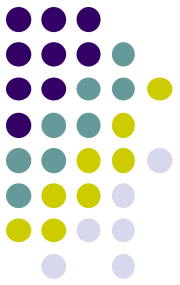
[O3] (Freestyle) Parallelize *any* program by OpenMP.

For more detail, please see [ppcomp-1-1](#) slides

Announcement of TSUBAME Maintenance



- **TSUBAME4 will be stopped** for maintenance
 - 9:00, Mon, Apr 22 – 17:00, Wed, Apr 24 (plan)
 - TSUBAME4 portal is also stopped
- **On Apr 22 class, we can not use TSUBAME**



Next Class:

- OpenMP(3)
 - Bottlenecks in parallel programs
 - Mutual exclusion, reduction