

Practical Parallel Computing (実践的並列コンピューティング)

Part 2: GPU

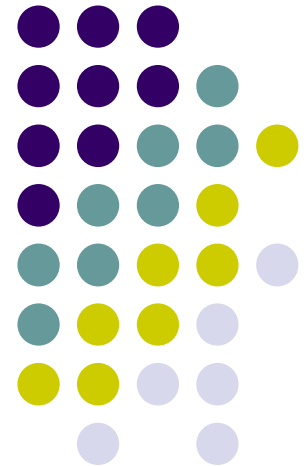
No 1: Overview and OpenACC

May 2, 2024

Toshio Endo

School of Computing & GSIC

endo@is.titech.ac.jp

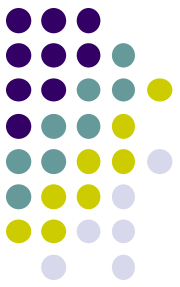




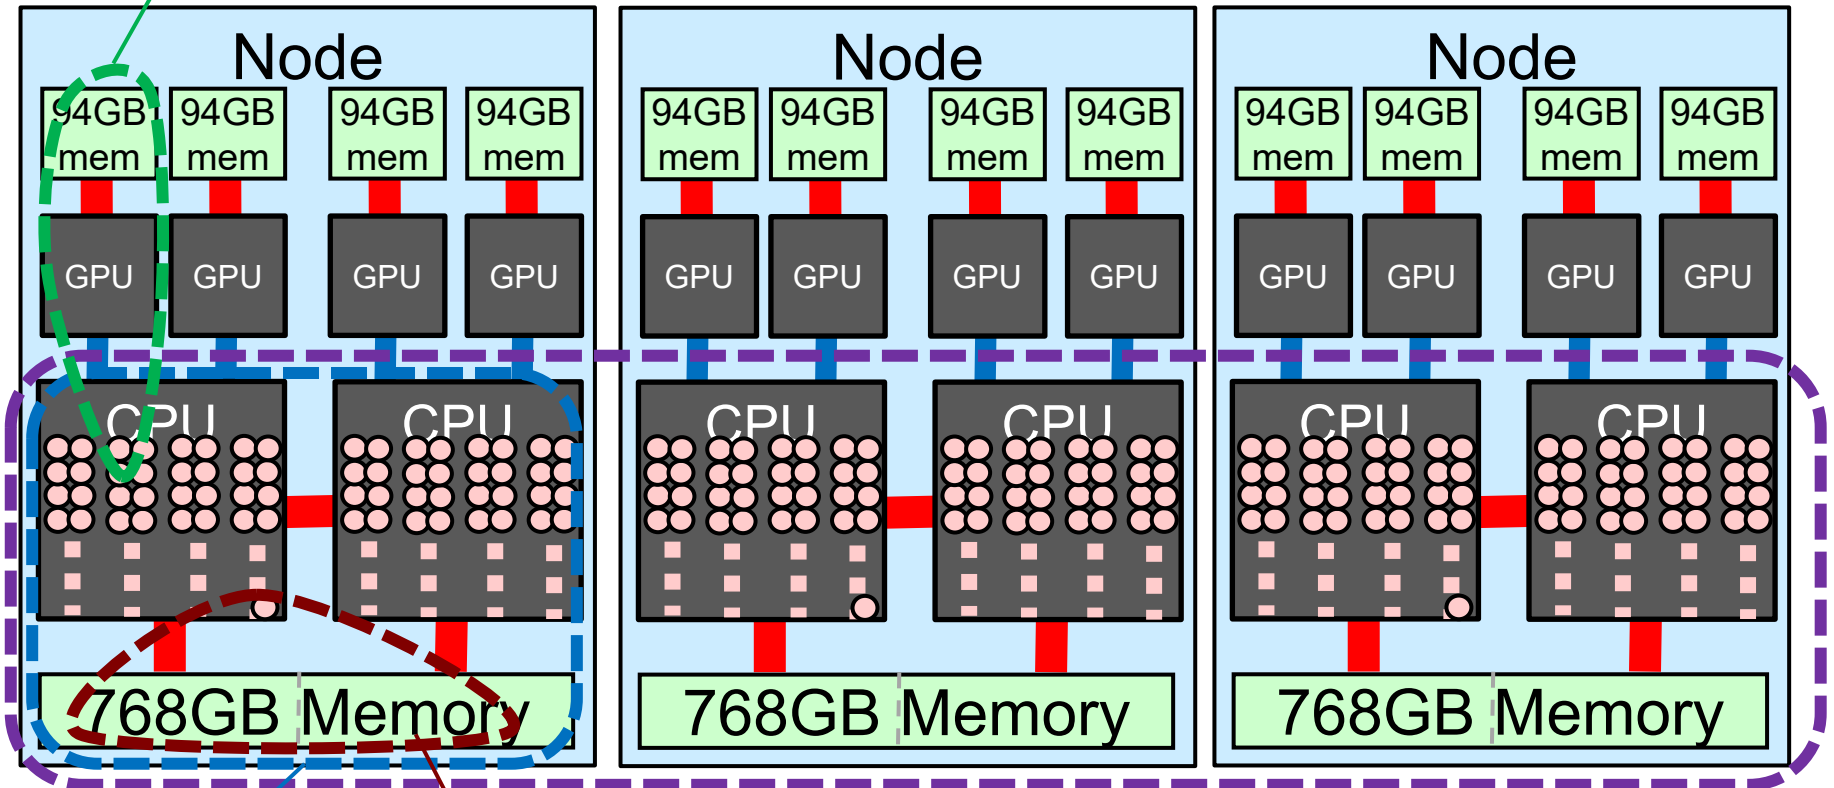
Overview of This Course

- Part 0: Introduction
 - 2 classes
- Part 1: OpenMP for shared memory programming
 - 4 classes
- Part 2: **GPU** programming
 - 4 classes **← We are here (1/4)**
 - OpenACC (1.5 classes) and CUDA (2.5 classes)
- Part 3: **MPI** for distributed memory programming
 - 3 classes

Parallel Programming Methods on TSUBAME




OpenACC/CUDA



MPI


OpenMP

Sequential




GPU Computing

- **Graphic processing units (GPU)** have been originally used for computing graphics (including video games)
- A high performance GPU has many cores
 - CPU: 2 to 32 cores. GPU: >1000 cores
 - The concept is called GPGPU (General-Purpose computing on GPU)
- GPGPU became popular since NVIDIA invented CUDA language in 2007
 - Recently it is popular for deep learning




TSUBAME4
node

GPU





A Compute Node with GPU



- A GPU has its distinct memory (**device memory**)
 - CPU memory is called **host memory**
- Many cores in a GPU share its device memory
- If there are multiple GPUs, each has its device memory



Characteristics of GPUs

A GPU is a board or a card attached to computers

→ It cannot work alone. Driven by CPUs


→ Different programming methods

Comparing EPYC 9654 (TSUBAME4's CPU) and

Tesla H100-94GB (TSUBAME4's GPU)

	1 CPU	1 GPU
Number of cores	96 cores (192 cores with 2CPUs)	16,896 CUDA cores (=128 x 132SMXs)
Clock Frequency	3.55GHz (with boost)	1.98GHz
Peak Computation Speed (double precision)	5.45TFlops	66.9TFlops
Memory Capacity	384GB (768GB shared by 2CPUs)	94GB

Notes on TSUBAME Interactive Node



In an interactive node, each user uses **0.5 GPU**

- 7,680 CUDA cores (128x60SMs), ~46GB memory

If you want use “a full GPU”, you may try “qsub” (described later)

Programming Environments for NVIDIA GPUs



- **CUDA** ← We will use after OpenACC
 - The most famous environment, designed by NVIDIA
 - C/Fortran + new syntaxes
 - Use “nvcc” command for compile
 - `module load cuda`
 - `nvcc ... XXX.cu`
 - For more general programs than OpenACC 😊
- **OpenACC** ← Today's topic
 - C/Fortran + directives (`#pragma acc ...`), Easier programming 😊
 - Supported by NVIDIA HPC SDK
 - `module load nvhpc`
 - `nvc -acc ... XXX.c`
 - For parallel programs with for-loops
- OpenMP 5, OpenCL...



An OpenACC Program Looks Like

C/C++/Fortran + directives

```
int a[100], b[100], c[100];
int i;
#pragma acc data copy(a,b,c)
#pragma acc kernels
#pragma acc loop independent
{
    for (i = 0; i < 100; i++) {
        a[i] = b[i]+c[i];
    }
}
```

Examples of **OpenACC**
directives

In this case, each directive has
an effect on the following
block/sentence

OpenACC is not so popular as OpenMP, unfortunately☹

- gcc's support is not so good
- On TSUBAME4, we use [NVIDIA HPC SDK](#)

OpenACC Version of “mm” sample




Available at </gs/bs/tga-ppcomp/24/mm-acc/>

A: a $(m \times k)$ matrix, B: a $(k \times n)$ matrix

C: a $(m \times n)$ matrix

$$C \leftarrow A \times B$$

- Algorithm with a triply-nested for-loop
- Supports variable matrix size.
 - Each matrix is expressed as a 1D array by *column-major* format
- Execution: `./mm [m] [n] [k]`





Using mm-acc Sample

[make sure that you are at a interactive node (rXn11)]

module load nvhpc *[Do once after login]*

cd ~/ppc24

cp -r /gs/bs/tga-ppcomp/24/**mm-acc** .

cd mm-acc

make

*[You will see some messages, and an executable file
“mm” is created]*

./mm 2000 2000 2000

Notes on Compiling OpenACC Programs



- NVIDIA HPC SDK on TSUBAME4.0
 - `module load nvhpc`, and then use `nvc` command
 - Use `-acc` option in compiling and linking
 - `-Minfo=accel` option outputs many information on parallelization

Example of output


```
:  
26, Generating copyin(A[:m*k])  
    Generating copy(C[:m*n])  
    Generating copyin(B[:k*n])  
    Loop is parallelizable  
:
```

They are not errors 😊

Kernel Region in OpenACC



```
int main()  
{  
  A:  
  #pragma acc kernels  
  {  
    B;  
  }  
  C;  
  #pragma acc kernels  
  {  
    D;  
    E;  
  }  
}
```




A sentence/block immediately after **#pragma acc kernels** is called a **kernel region**, executed on GPU

- We don't need to specify number of threads (it is hard to specify explicitly)

Data Movement between CPU and GPU



- We need to move data between CPU and GPU
 - Host (CPU) memory and Device (GPU) memory are distinct, like distributed memory
 - Threads on a GPU share the device memory




For this purpose, we use `#pragma acc data` directive
→ This defines a `data region`

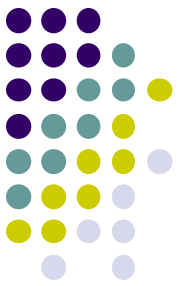
Data Directives to use GPU memory



```
int main()
{
    A;
    #pragma acc data copy(x,y)
    {
        #pragma acc kernels
        {
            B;
        }
        C;
        #pragma acc kernels
        {
            D;
        }
    }
    E;
}
```



- Data region may contain 1 or more kernel regions
- Data movement occurs at beginning and end of data region



Data Directive (1)

- Arrays (like a):
 - we can write array names if the sizes are statically declared → entire array is copied
- Pointers as arrays (like b):

cf) `b[0:20]`

start index number of elements

 - Partial copying like `b[10:5]` or `a[4:4]` work
- Scalar variables (like x):
 - You can omit `copy(x)` → The compiler detects automatically 😊

```
int x;  
float a[10];  
double *b = (double*)  
    malloc(20*sizeof(double));  
:  
#pragma acc data copy(x, a, b[0:20])  
:
```

Same meaning

```
#pragma acc data copy(a[0:10], b[0:20])
```




Data Directive (2)

- Directions of copying
 - ... `data copyin(...)`: Copy CPU→GPU at the begininng
 - ... `data copyout(...)`: Copy GPU→CPU at the end
 - ... `data copy(...)`: Do both

Optimization of data movement will help speedup



Loop Directive

```
int a[100], b[100], c[100];  
int i;  
#pragma acc data copy(a,b,c)  
#pragma acc kernels  
#pragma acc loop independent  
for (i = 0; i < 100; i++) {  
    a[i] = b[i]+c[i];  
}
```

- ... **loop independent**: Iterations are done in parallel by multiple GPU threads
- ... **loop seq**: Done sequentially. Not be parallelized
- ... **loop**: Compiler decides

- **#pragma acc loop** must be included in “**acc kernels**” or “acc parallel”
- Directly followed by “for” loop
 - The loop must have a loop counter, as in OpenMP
 - List/tree traversal is NG

OpenACC Version of mm (mm-acc/mm.c)



```
#pragma acc data copyin(A[0:m*k],B[0:k*n]),copy(C[0:m*n])
```

```
#pragma acc kernels
```

```
#pragma acc loop independent
```

```
    for (j = 0; j < n; j++) {
```

```
#pragma acc loop seq
```

```
        for (l = 0; l < k; l++) {
```

```
#pragma acc loop independent
```

```
            for (i = 0; i < m; i++) {
```

```
                Ci,j += Ai,l * Bl,j;
```

```
            } } }
```

← We can omit GPU → CPU copy of A,B

← For each column in C

← For dot product

← For each row in C

- Each element in C can be computed in parallel (i-loop, j-loop)
- Computation of a single C element is sequential (l-loop)

Submitting a GPU Job to the Job Scheduler



- Sequential version
 - see [mm](#) directory

- OpenACC version
 - see [mm-acc](#) directory
 - To use a GPU, use **node_q** type
 - (node_h or node_f types for multi-GPU)

mm/job.sh

resource type
and count

maximum
run time

```
#!/bin/sh
#$ -cwd
#$ -l cpu_4=1
#$ -l h_rt=00:10:00

./mm 2000 2000 2000
```

mm-acc/job.sh

```
#!/bin/sh
#$ -cwd
#$ -l node_q=1
#$ -l h_rt=00:10:00

./mm 2000 2000 2000
```

- Job submission
 - `qsub job.sh`



0.5 GPU vs 1 GPU

$m=n=k=2000$

Interactive node (0.5GPU)

node_q (1GPU)

We can ignore speed of the first computation

Matmul took 363154 us --> 44.058 GFlops

Matmul took 35186 us --> 454.726 GFlops

Matmul took 35156 us --> 455.114 GFlops

Matmul took 35181 us --> 454.791 GFlops

Matmul took 35130 us --> 455.451 GFlops

Matmul took 1816737 us --> 8.807 GFlops

Matmul took 20691 us --> 773.283 GFlops

Matmul took 20712 us --> 772.499 GFlops

Matmul took 20577 us --> 777.567 GFlops

Matmul took 20642 us --> 775.119 GFlops

About 1.7 times speedup


The first computation suffers slowdown by initialization of GPU etc.
The similar slowdown also occurs in CPU programs

OpenACC version of “pi” sample



Estimate approximation of π (circumference/diameter) by approximation of integration

- Available at </gs/bs/tga-ppcomp/24/pi-acc/>
- Method
 - Let SUM be approximation of the yellow area
 - $4 \times \text{SUM} \rightarrow \pi$
- Execution: `./pi [n]`
 - n: Number of division
 - Cf) `./pi 100000000`
- Compute complexity: $O(n)$



$$dx = 1/n$$
$$y = \sqrt{1-x*x}$$

Algorithm of “pi”



OpenMP

```
double pi(int n) {  
    int i;  
    double sum = 0.0;  
    double dx = 1.0 / (double)n;  
  
    #pragma omp parallel  
    #pragma omp for reduction(+:sum)  
    for (i = 0; i < n; i++) {  
        double x = (double)i * dx;  
        double y = sqrt(1.0 - x*x);  
        sum += dx*y;  
    }  
  
    return 4.0*sum; }
```

OpenACC

```
double pi(int n) {  
    int i;  
    double sum = 0.0;  
    double dx = 1.0 / (double)n;  
  
    #pragma acc kernels  
    #pragma acc loop independent reduction(+:sum)  
    for (i = 0; i < n; i++) {  
        double x = (double)i * dx;  
        double y = sqrt(1.0 - x*x);  
        sum += dx*y;  
    }  
  
    return 4.0*sum; }
```

✖ For scalar variables, “data copy”
is omitted




Notes on Number of Threads

- In OpenMP, the number of threads is set by `OMP_NUM_THREADS`
- In OpenACC, the number is **automatically** determined per loop
- In OpenMP, thread ID is obtained by `omp_get_thread_num()`
- In OpenACC, we **cannot see thread ID**

Assignments in this Course



- There is homework for each part. Submissions of reports for **2 parts** are required





Assignments in GPU Part (1)

Choose one of [G1]—[G3], and submit a report

Due date: May 30 (Thursday)

[G1] Parallelize “diffusion” sample program by OpenACC or CUDA

- You can use Makefile in </gs/bs/tga-ppcomp/24/diffusion-acc/> or </gs/bs/tga-ppcomp/24/diffusion-cuda/>

Optional:

- To make array sizes variable parameters
- To compare OpenACC vs CUDA
- To improve performance further
 - Different assignment of threads and elements (CUDA), etc

Assignments in GPU Part(2)



[G2] Evaluate speed of “mm-acc” or “mm-cuda” in detail

mm-acc: </gs/bs/tga-ppcomp/24/mm-acc/>

mm-cuda: </gs/bs/tga-ppcomp/24/mm-cuda/>

- Use various matrices sizes
- Evaluate effects of data transfer cost
- Compare with CPU (OpenMP) version

Optional:

- To use different loop orders
- To evaluate both mm-acc and mm-cuda
- To change/improve the program
 - Different assignment of threads and elements (CUDA) etc



Assignments in GPU Part (3)

[G3] (Freestyle) Parallelize *any* program by OpenACC or CUDA.

- cf) A problem related to your research
- “sort” sample on GPU?
 - The quick sort may be hard on GPU (There is no “task” syntax)
 - → Bitonic sort?
- More challenging one for parallelization is better
 - cf) Partial computations have dependency with each other

Notes in Report Submission (1)



- Submit the followings via **T2SCHOLA**
 - (1) **A report document**
 - PDF, MS-Word or text file
 - 2 pages or more
 - in English or Japanese (日本語もok)
 - (2) **Source code files** of your program
 - Try “zip” to submit multiple files

Notes in Report Submission (2)



The report document should include:

- Which problem you have chosen
 - In Part 2, describe which you used, **OpenACC** or **CUDA**?
- How you parallelized
 - It is even better if you mention efforts for high performance or new functions
- Performance evaluation on TSUBAME
 - With varying number of threads
 - With varying problem sizes
 - Discussion with your findings
 - Other machines than TSUBAME are ok, if available



Next Class:

- GPU Programming (2)
 - OpenACC
 - Improving data copy
 - Improving loop parallelization
 - Introduction of CUDA
- Schedule
 - Mon, May 6: No classes (national holiday)
 - Thu, May 9: GPU (2)
 - Mon, May 13: GPU (3)
 - Thu, May 16: No classes (cancelled/休講)